


# Practices in visual computing 2

## Lab3: GANs

Simon Fraser University  
Spring 2025

# Generative learning



# Generative learning

## Explicit Models:

- train through directly optimizing  $p(x)$
- Through Maximum likelihood estimation (MLE)
- Example: VAEs

## Implicit Models:

- Train to sample from  $p(x)$
- Example: GANs

# GAN Conceptual Introduction

Two networks (generator G and discriminator D) compete in a **minimax** game.

**Generator (G)**: Learns to produce “fake” data that mimic real data.


**Discriminator (D)**: Learns to distinguish between real and fake data.

As training progresses: G gets better at fooling D.


D gets better at identifying fakes.

Goal: The generator produces samples indistinguishable from real ones, and the discriminator fails to tell them apart.

# Training GANs



# GAN Architecture Diagram



# Training GANs

Jointly Train Generator and Discriminator through a minmax loss

- Alternate between:
  - Gradient ascent for Discriminator

$$\max_{\theta_d} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p_z} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$

- Gradient ascent for Discriminator

$$\min [\mathbb{E}_{z \sim p_z} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$

# Training GANs


- Gradient ascent for Discriminator

$$\max_{\theta_d} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p_z} \log(1 - D_{\theta_d}(G_{\theta_g}(x)))]$$

- Gradient ascent for Discriminator

$$\min [\mathbb{E}_{z \sim p_z} \log(1 - D_{\theta_d}(G_{\theta_g}(x)))]$$

Gradient small when  
discriminator is winning





# Training GANs

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right).$$


**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---


# GANs: Convolutional architecture

Generator is a upsampling network with fractionally-strided Convolutions.



# GANs: Convolutional architecture


Generator is a upsampling network with **fractionally-strided** Convolutions.



$$o = (i - 1) \times s + k - 2p$$


# GANs: Convolutional architecture

Generator is a upsampling network with fractionally-strided Convolutions.



# GANs: Convolutional architecture

- Generator is a upsampling network with fractionally-strided Convolutions.
- Discriminator is a CNN.



# Tips and Tricks

Training GANs is difficult

- Stability between Generator and Discriminator
  - Vanishing Gradients - Exploding Gradients
- Mode Collapse
  - The generator repeatedly produces similar or identical samples.

# Tips and Tricks

There are general tips that can help train GANs

From <https://github.com/soumith/ganhacks>

Normalize input in range  $(-1, 1)$


- Use Tanh activation for the last layer of the Generator output
- 1. Avoid sparse gradients
  - Avoid using ReLU and maxpool
- 2. Track failures early
  - If Discriminator gradients go to zero: failure mode
  - If there is high variance/spikes in the Discriminator loss: failure mode
  - If loss of Generator steadily decreases it is fooling the Discriminator with garbage
- 3. Etc.

# GAN progress on face generation





# GAN MNIST



# Reference

[https://developers.google.com/machine-learning/gan/gan\\_structure](https://developers.google.com/machine-learning/gan/gan_structure)

<https://machinelearningmastery.com/how-to-code-the-generative-adversarial-network-training-algorithm-and-loss-functions/>

<https://towardsdatascience.com/gan-ways-to-improve-gan-performance-acf37f9f59b>

Slides from last year of CMPT743